

How K was chosen — the K in use per cell (K_selection.csv), reduced model on every row

fill and number = the K in use (selected_K), grid K = 1..10 · sparkline = that cell's own BIC-vs-K curve, red dot = its minimum

'by hand' = selected_K differs from the default (first K whose BIC gain < 5% of the total drop) · 'n/10 conv' = the other K did not converge

'N var/M sh' = the random part asks for N variances from M shoulders, so K is NOT identifiable: it describes the flexibility the optimiser bore, not the trajectory

